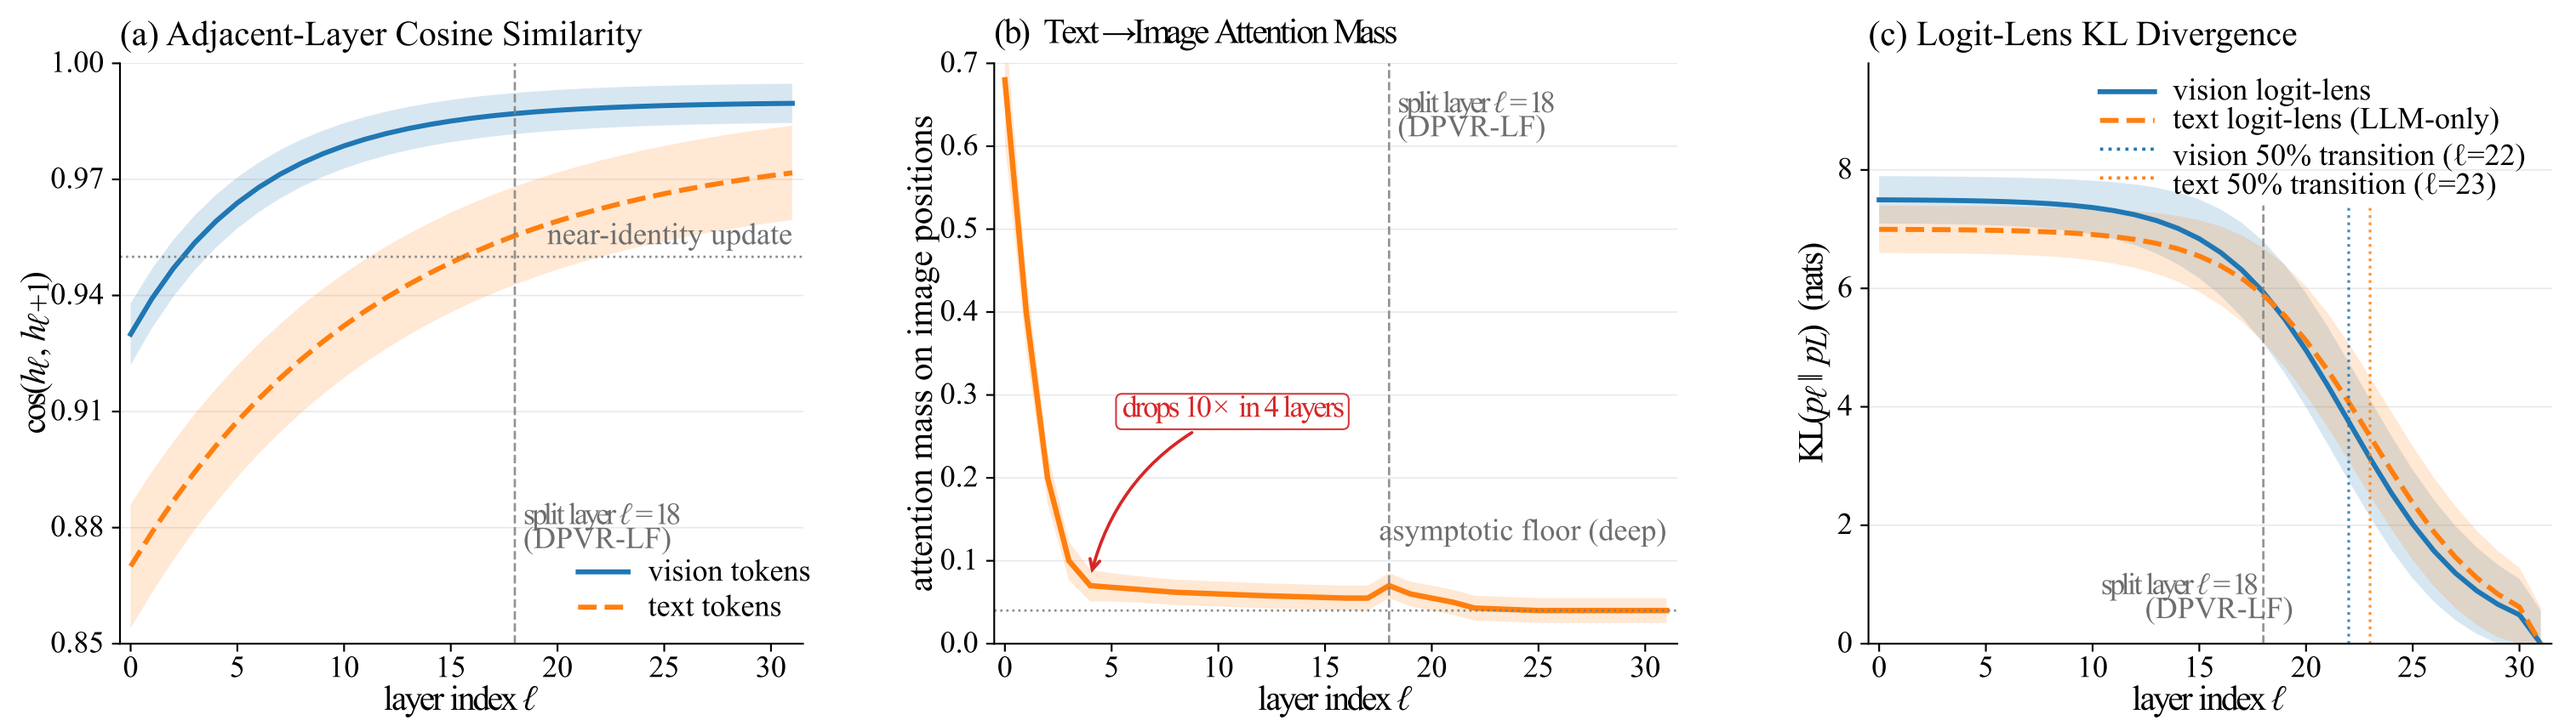

(a) Adjacent-Layer Cosine Similarity
cos(hℓ, hℓ+1)
layer index ℓ
near-identity update
split layer ℓ = 18
(DPVR-LF)
vision tokens
text tokens
(b) Text→Image Attention Mass
attention mass on image positions
layer index ℓ
split layer ℓ = 18
(DPVR-LF)
drops 10× in 4 layers
asymptotic floor (deep)
(c) Logit-Lens KL Divergence
KL(pℓ ‖ pL) (nats)
layer index ℓ
vision logit-lens
text logit-lens (LLM-only)
vision 50% transition (ℓ=22)
text 50% transition (ℓ=23)
split layer ℓ = 18
(DPVR-LF)